

# CIS664-Knowledge Discovery and Data Mining

## Mining Association Rules

Vasileios Megalooikonomou  
Dept. of Computer and Information Sciences  
Temple University

(based on notes by Jiawei Han and Micheline Kamber)

# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

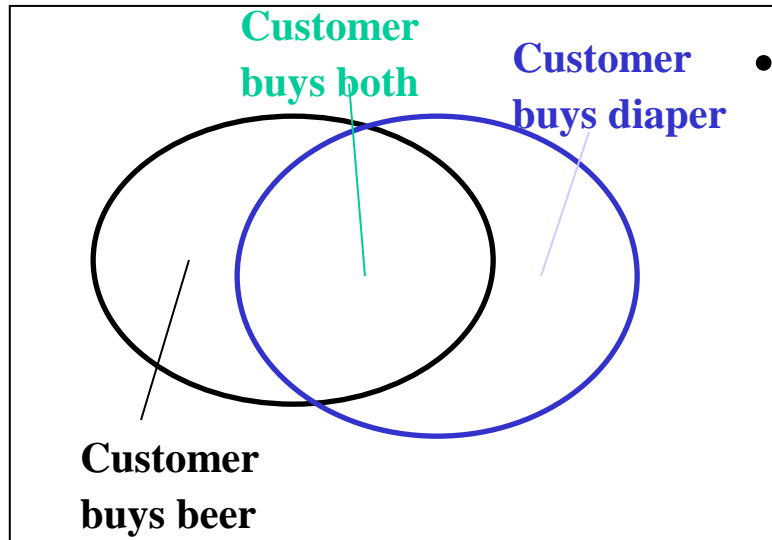
# Association Mining?

- Association rule mining:
  - Finding frequent patterns, associations, correlations, or causal structures among sets of items or objects in transaction databases, relational databases, and other information repositories.
- Applications:
  - Basket data analysis, cross-marketing, catalog design, loss-leader analysis, clustering, classification, etc.
- Examples.
  - Rule form: “**Body**  $\rightarrow$  **Head** [support, confidence]”.
  - $\text{buys}(x, \text{“diapers”}) \rightarrow \text{buys}(x, \text{“beers”}) [0.5\%, 60\%]$
  - $\text{major}(x, \text{“CS”}) \wedge \text{takes}(x, \text{“DB”}) \rightarrow \text{grade}(x, \text{“A”}) [1\%, 75\%]$

# Association Rules: Basic Concepts

- Given: (1) database of transactions, (2) each transaction is a list of items (purchased by a customer in a visit)
- Find: all rules that correlate the presence of one set of items with that of another set of items
  - E.g., *98% of people who purchase tires and auto accessories also get automotive services done*
- Applications
  - $* \Rightarrow \textit{Maintenance Agreement}$  (What the store should do to boost Maintenance Agreement sales)
  - *Home Electronics*  $\Rightarrow *$  (What other products should the store stocks up?)
  - Attached mailing in direct marketing
  - Detecting “ping-pong”ing of patients, faulty “collisions”

# Interestingness Measures: Support and Confidence



- Find all the rules  $X \& Y \Rightarrow Z$  with minimum confidence and support
  - **support,  $s$** , probability that a transaction contains  $\{X \& Y \& Z\}$
  - **confidence,  $c$** , conditional probability that a transaction having  $\{X \& Y\}$  also contains  $Z$

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

*Let minimum support 50%, and minimum confidence 50%, we have*

- $A \Rightarrow C$  (50%, 66.6%)
- $C \Rightarrow A$  (50%, 100%)

# Association Rule Mining: A Road Map

- Boolean vs. quantitative associations (Based on the types of values handled)
  - $\text{buys}(x, \text{"SQLServer"}) \wedge \text{buys}(x, \text{"DMBook"}) \rightarrow \text{buys}(x, \text{"DBMiner"})$   
[0.2%, 60%]
  - $\text{age}(x, \text{"30..39"}) \wedge \text{income}(x, \text{"42..48K"}) \rightarrow \text{buys}(x, \text{"PC"})$  [1%, 75%]
- Single dimension vs. multiple dimensional associations (each distinct predicate of a rule is a dimension)
- Single level vs. multiple-level analysis (consider multiple levels of abstraction)
  - What brands of beers are associated with what brands of diapers?
- Extensions
  - Correlation, causality analysis
    - Association does not necessarily imply correlation or causality
  - Maxpatterns (a frequent pattern s.t. any proper subpattern is not frequent) and closed itemsets (if there exist no proper superset  $c'$  of  $c$  s.t. any transaction containing  $c$  also contains  $c'$ )

# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

# Mining Association Rules—An Example

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Min. support 50%  
Min. confidence 50%

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

For rule  $A \Rightarrow C$ :

support = support( $\{A \cup C\}$ ) = 50%

confidence = support( $\{A \cup C\}$ )/support( $\{A\}$ ) = 66.6%

The **Apriori** principle:

Any subset of a frequent itemset must be frequent



# Mining Frequent Itemsets

- Find the *frequent itemsets*: the sets of items that have minimum support
  - A subset of a frequent itemset must also be a frequent itemset
    - i.e., if  $\{AB\}$  is a frequent itemset, both  $\{A\}$  and  $\{B\}$  should be a frequent itemset
  - Iteratively find frequent itemsets with cardinality from 1 to  $k$  ( $k$ -itemset)
- Use the frequent itemsets to generate association rules.

# The Apriori Algorithm: Basic idea

- **Join Step:**  $C_k$  is generated by joining  $L_{k-1}$  with itself
- **Prune Step:** Any  $(k-1)$ -itemset that is not frequent cannot be a subset of a frequent  $k$ -itemset

- Pseudo-code:

$C_k$ : Candidate itemset of size  $k$

$L_k$ : frequent itemset of size  $k$

$L_1 = \{\text{frequent items}\};$

**for** ( $k = 1; L_k \neq \emptyset; k++$ ) **do begin**

$C_{k+1}$  = candidates generated from  $L_k$ ;

**for each** transaction  $t$  in database **do**

increment the count of all candidates in  $C_{k+1}$   
that are contained in  $t$

$L_{k+1}$  = candidates in  $C_{k+1}$  with min\_support

**end**

**return**  $\cup_k L_k$ ;

# The Apriori Algorithm — Example

Database D

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

$C_1$

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

Scan D

$L_1$

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

$C_2$

itemset	sup
{1 2}	1
{1 3}	2
{1 5}	1
{2 3}	2
{2 5}	3
{3 5}	2

Scan D

$C_2$

itemset
{1 2}
{1 3}
{1 5}
{2 3}
{2 5}
{3 5}

$L_2$

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2

$C_3$

itemset
{2 3 5}

$L_3$

itemset	sup
{2 3 5}	2

# How to Generate Candidates?

- Suppose the items in  $L_{k-1}$  are listed in an order
- Step 1: self-joining  $L_{k-1}$ 
  - insert into  $C_k$
  - select  $p.item_1, p.item_2, ..., p.item_{k-1}, q.item_{k-1}$
  - from  $L_{k-1} p, L_{k-1} q$
  - where  $p.item_1=q.item_1, ..., p.item_{k-2}=q.item_{k-2}, p.item_{k-1} < q.item_{k-1}$
- Step 2: pruning
  - forall *itemsets*  $c$  in  $C_k$  do
    - forall  $(k-1)$ -subsets  $s$  of  $c$  do
      - if ( $s$  is not in  $L_{k-1}$ ) then delete  $c$  from  $C_k$

# How to Count Supports of Candidates?

- Why is counting supports of candidates a problem?
  - The total number of candidates can be huge
  - Each transaction may contain many candidates
- Method:
  - Candidate itemsets are stored in a *hash-tree*
  - *Leafnode* of hash-tree contains a list of itemsets and counts
  - *Interior node* contains a hash table
  - *Subset function*: finds all the candidates contained in a transaction

# Example of Generating Candidates

- $L_3 = \{abc, abd, acd, ace, bcd\}$
- Self-joining:  $L_3 * L_3$ 
  - $abcd$  from  $abc$  and  $abd$
  - $acde$  from  $acd$  and  $ace$
- Pruning:
  - $acde$  is removed because  $ade$  is not in  $L_3$
- $C_4 = \{abcd\}$

# Improving Apriori's Efficiency

- **Hash-based itemset counting:** A  $k$ -itemset whose corresponding hashing bucket count is below the threshold cannot be frequent
- **Transaction reduction:** A transaction that does not contain any frequent  $k$ -itemset is useless in subsequent scans
- **Partitioning:** Any itemset that is potentially frequent in DB must be frequent in at least one of the partitions of DB
- **Sampling:** mining on a subset of given data, need a lower support threshold + a method to determine the completeness
- **Dynamic itemset counting:** add new candidate itemsets immediately (unlike Apriori) when all of their subsets are estimated to be frequent

# Is Apriori Fast Enough? — Performance Bottlenecks

- The core of the Apriori algorithm:
  - Use frequent  $(k-1)$ -itemsets to generate candidate frequent  $k$ -itemsets
  - Use database scan and pattern matching to collect counts for the candidate itemsets
- The bottleneck of *Apriori*: candidate generation
  - Huge candidate sets:
    - $10^4$  frequent 1-itemset will generate  $10^7$  candidate 2-itemsets
    - To discover a frequent pattern of size 100, e.g.,  $\{a_1, a_2, \dots, a_{100}\}$ , one needs to generate  $2^{100} \approx 10^{30}$  candidates.
  - Multiple scans of database:
    - Needs  $(n+1)$  scans,  $n$  is the length of the longest pattern



# Mining Frequent Patterns Without Candidate Generation

- Compress a large database into a compact, Frequent-  
Pattern tree (FP-tree) structure
  - highly condensed, but complete for frequent pattern mining
  - avoid costly database scans
- Develop an efficient, FP-tree-based frequent pattern mining method
  - A divide-and-conquer methodology: decompose mining tasks into smaller ones
  - Avoid candidate generation: sub-database test only!

# Construct FP-tree from a Transaction DB

<i>TID</i>	<i>Items bought</i>	<i>(ordered) frequent items</i>
100	{ <i>f, a, c, d, g, i, m, p</i> }	{ <i>f, c, a, m, p</i> }
200	<i>b, m</i> { <i>a, b, c, f, l, m, o</i> }	{ <i>f, c, a</i> ,
300	{ <i>b, f, h, j, o</i> }	{ <i>f, b</i> }
400	{ <i>b, c, k, s, p</i> }	{ <i>c, b, p</i> }
500	{ <i>a, f, c, e, l, p, m, n</i> }	{ <i>f, c, a, m, p</i> }

*min\_support* = 0.5

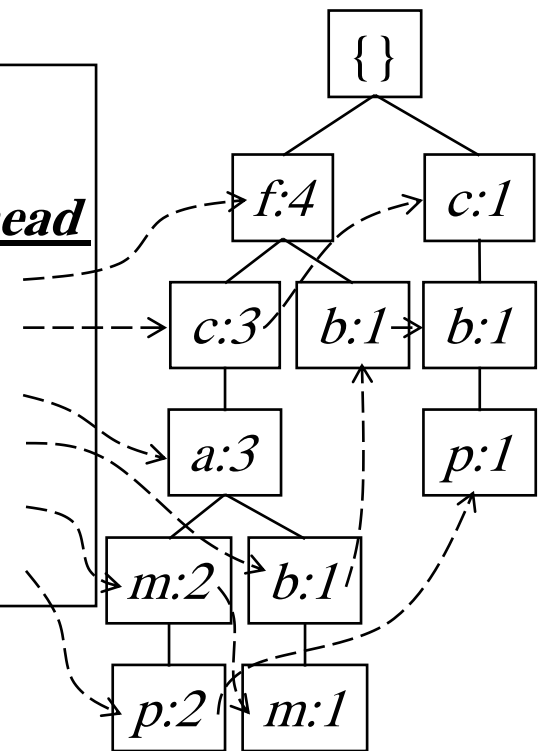
Steps:

1. Scan DB once, find frequent 1-itemset (single item pattern)
2. Order frequent items in frequency descending order
3. Scan DB again, construct FP-tree

**Header Table**

***Item frequency head***

<i>f</i>	4
<i>c</i>	4
<i>a</i>	3
<i>b</i>	3
<i>m</i>	3
<i>p</i>	3



# Benefits of the FP-tree Structure

- Completeness:
  - never breaks a long pattern of any transaction
  - preserves complete information for frequent pattern mining
- Compactness
  - reduce irrelevant information—infrequent items are gone
  - frequency descending ordering: more frequent items are more likely to be shared
  - never be larger than the original database (if not count node-links and counts)

# Mining Frequent Patterns Using FP-tree


- General idea (divide-and-conquer)
  - Recursively grow frequent pattern path using the FP-tree
- Method
  - For each item, construct its **conditional pattern-base**, and then its **conditional FP-tree**
  - Repeat the process on each newly created conditional FP-tree
  - Until the resulting FP-tree is **empty**, or it contains **only one path** (single path will generate all the combinations of its sub-paths, each of which is a frequent pattern)

# Major Steps to Mine FP-tree

- 1) Construct conditional pattern base for each node in the FP-tree
- 2) Construct conditional FP-tree from each conditional pattern-base
- 3) Recursively mine conditional FP-trees and grow frequent patterns obtained so far
  - If the conditional FP-tree contains a single path, simply enumerate all the patterns

# Step 1: From FP-tree to Conditional Pattern Base

- Starting at the frequent header table in the FP-tree
- Traverse the FP-tree by following the link of each frequent item
- Accumulate all of transformed prefix paths of that item to form a conditional pattern base



**Conditional pattern bases**


<u>item</u>	<u>cond. pattern base</u>
<i>c</i>	<i>f:3</i>
<i>a</i>	<i>fc:3</i>
<i>b</i>	<i>fca:1, f:1, c:1</i>
<i>m</i>	<i>fca:2, fcab:1</i>
<i>p</i>	<i>fcam:2, cb:1</i>

# Properties of FP-tree for Conditional Pattern Base Construction

- Node-link property
  - For any frequent item  $a_i$ , all the possible frequent patterns that contain  $a_i$  can be obtained by following  $a_i$ 's node-links, starting from  $a_i$ 's head in the FP-tree header
- Prefix path property
  - To calculate the frequent patterns for a node  $a_i$  in a path  $P$ , only the prefix sub-path of  $a_i$  in  $P$  need to be accumulated, and its frequency count should carry the same count as node  $a_i$ .

# Step 2: Construct Conditional FP-tree

- For each pattern-base
  - Accumulate the count for each item in the base
  - Construct the FP-tree for the frequent items of the pattern base



*m*-conditional  
pattern base:

*fca:2, fcab:1*



{}  
|  
*f:3*  
|  
*c:3*  
|  
*a:3*

All frequent patterns  
concerning *m*

*m*,  
*fm, cm, am*,  
*fcm, fam, cam*,  
*fcam*


*m*-conditional FP-tree



# Mining Frequent Patterns by Creating Conditional Pattern-Bases

Item	Conditional pattern-base	Conditional FP-tree
p	{(fcam:2), (cb:1)}	{(c:3)} p
m	{(fca:2), (fcab:1)}	{(f:3, c:3, a:3)} m
b	{(fca:1), (f:1), (c:1)}	Empty
a	{(fc:3)}	{(f:3, c:3)} a
c	{(f:3)}	{(f:3)} c
f	Empty	Empty

# Step 3: Recursively mine the conditional FP-tree



# Single FP-tree Path Generation

- Suppose an FP-tree  $T$  has a single path  $P$
- The complete set of frequent pattern of  $T$  can be generated by enumeration of all the combinations of the sub-paths of  $P$

$\{\}$   
|  
 $f:3$   
|  
 $c:3$   
|  
 $a:3$



All frequent patterns  
concerning  $m$

$m$ ,  
 $fm, cm, am$ ,  
 $fcm, fam, cam$ ,  
 $fcam$

*$m$ -conditional FP-tree*


# Principles of Frequent Pattern Growth

- Pattern growth property
  - Let  $\alpha$  be a frequent itemset in DB,  $B$  be  $\alpha$ 's conditional pattern base, and  $\beta$  be an itemset in  $B$ . Then  $\alpha \cup \beta$  is a frequent itemset in DB iff  $\beta$  is frequent in  $B$ .
- “*abcdef*” is a frequent pattern, if and only if
  - “*abcde*” is a frequent pattern, and
  - “*f*” is frequent in the set of transactions containing “*abcde*”

# Why Is Frequent Pattern Growth Fast?


- Our performance study shows
  - FP-growth is an order of magnitude faster than Apriori, and is also faster than tree-projection
- Reasoning
  - No candidate generation, no candidate test
  - Use compact data structure
  - Eliminate repeated database scan
  - Basic operation is counting and FP-tree building

# FP-growth vs. Apriori: Scalability With the Support Threshold



# FP-growth vs. Tree-Projection: Scalability with Support Threshold

Data set T25I20D100K

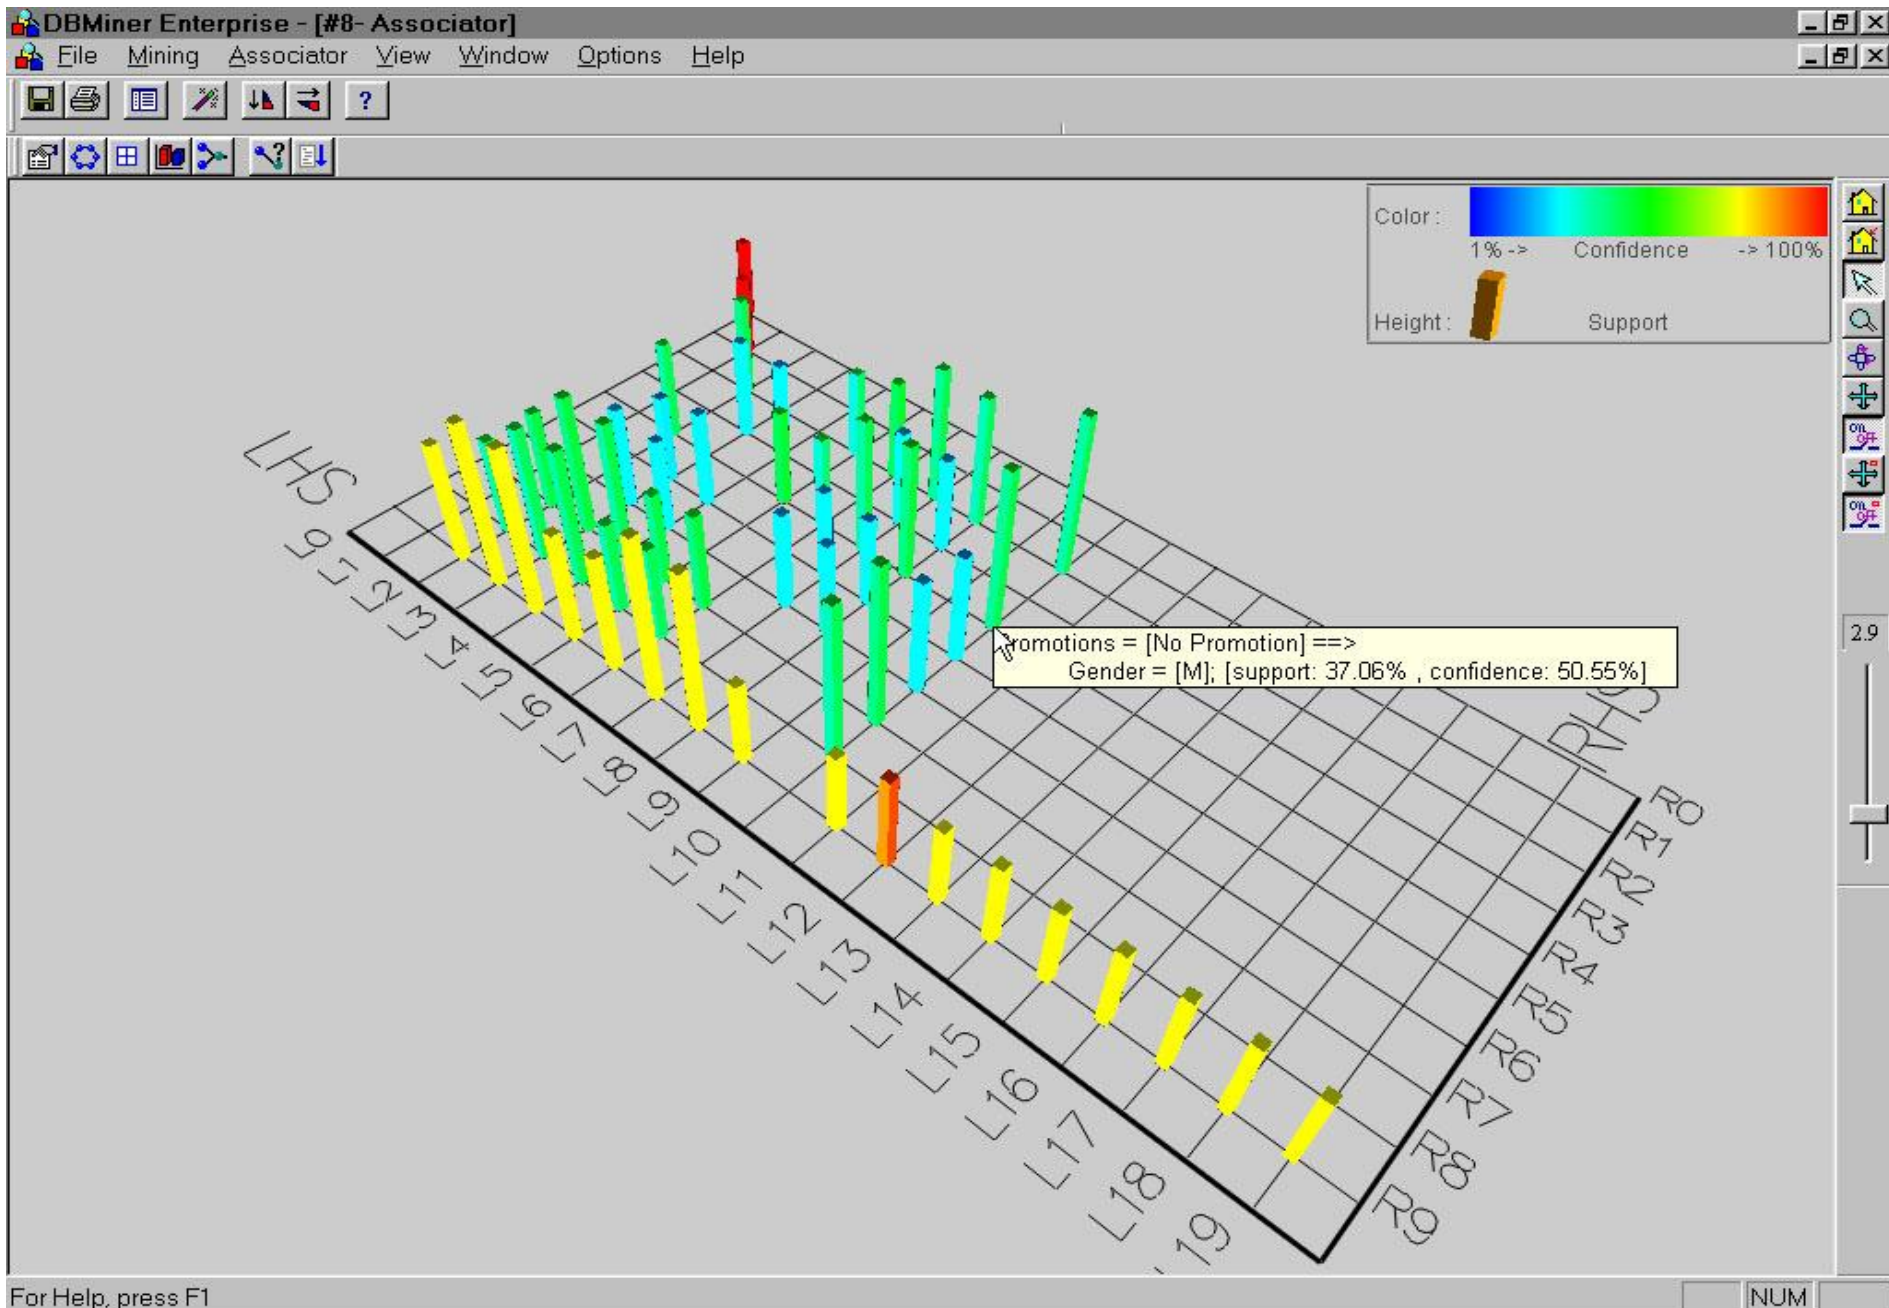


# Presentation of Association Rules (Table Form )

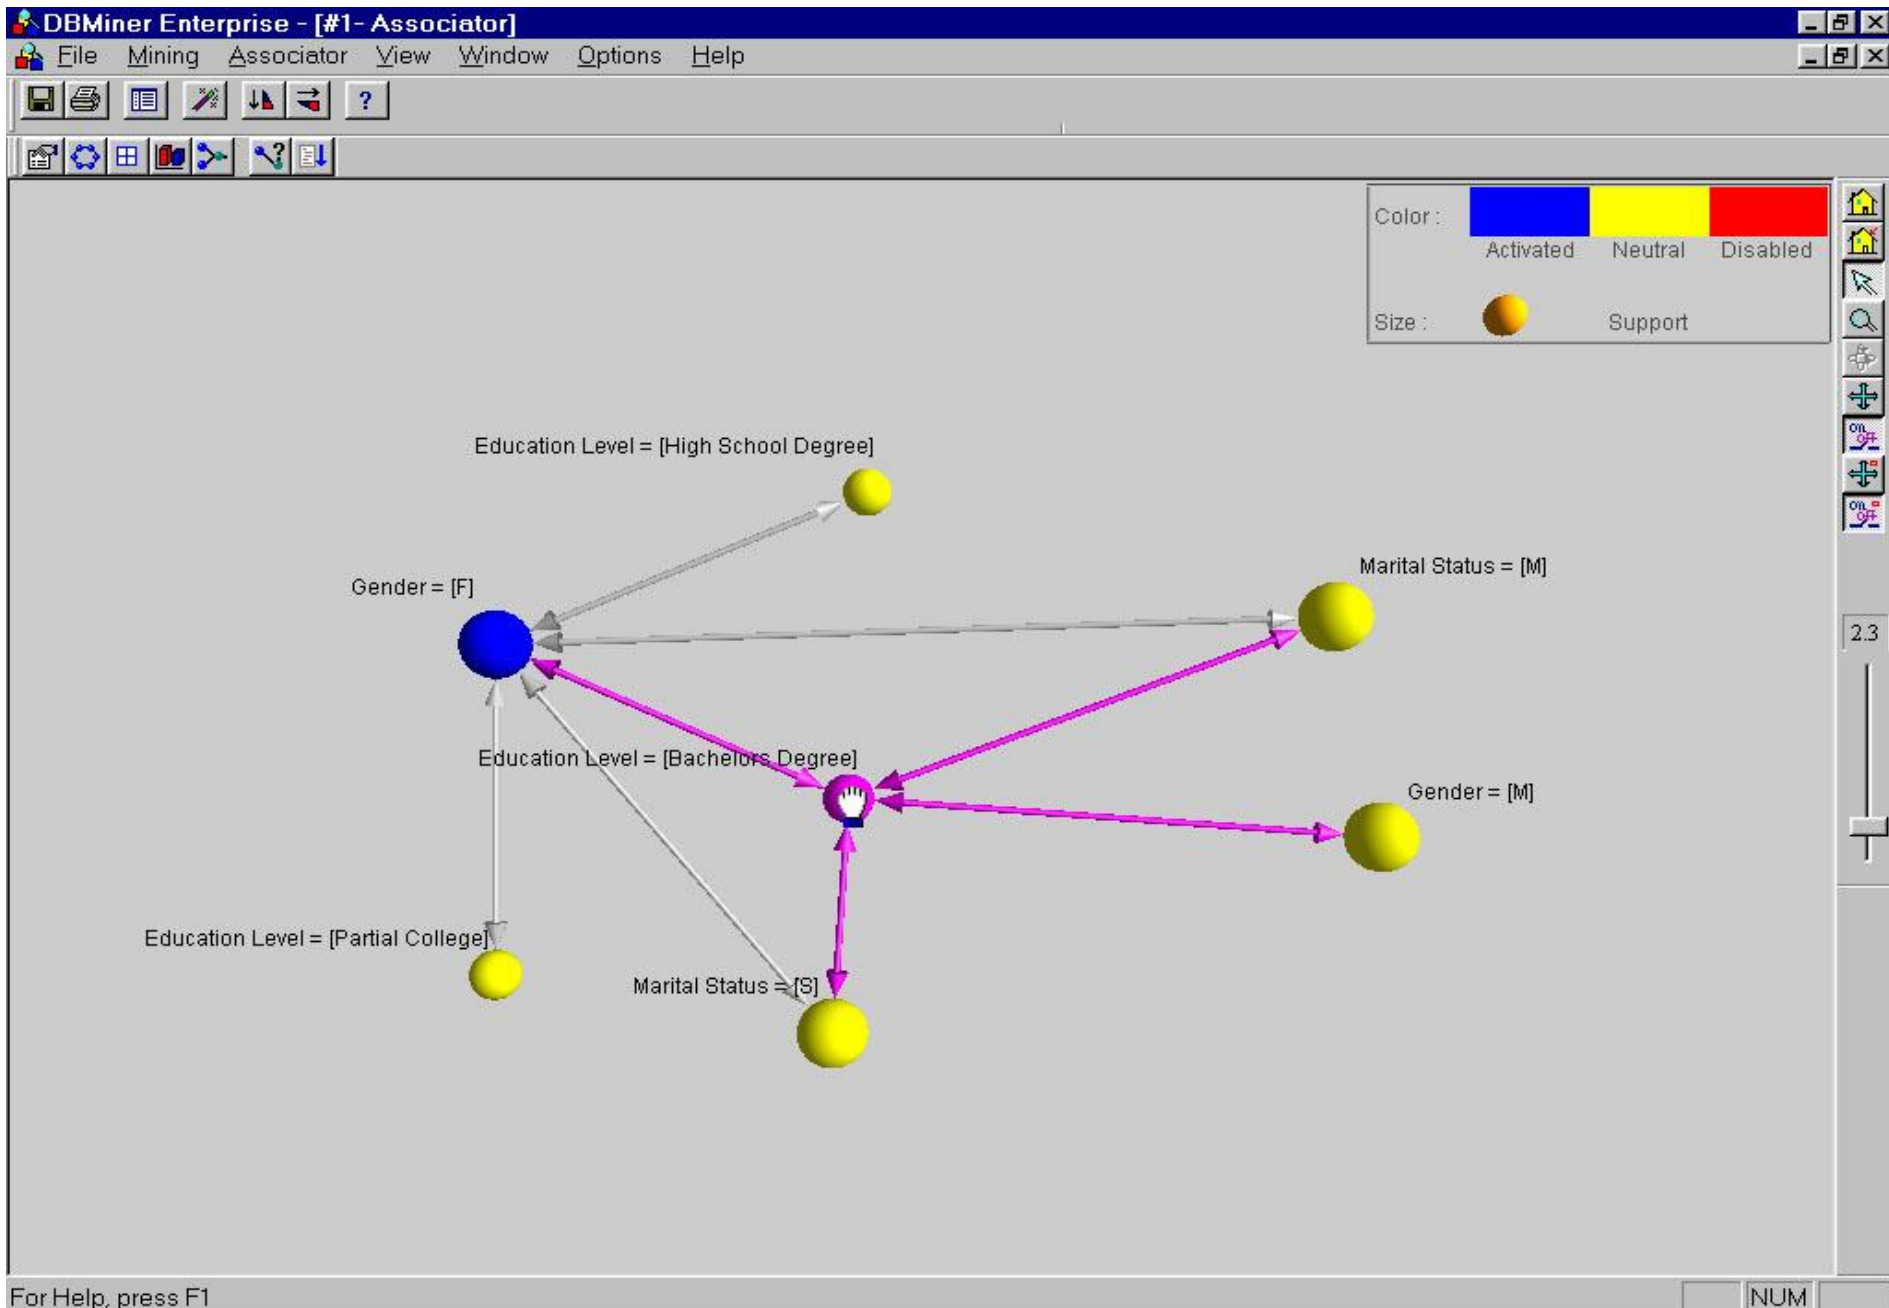
	Body	Implies	Head	Supp (%)	Conf (%)	F	G	H	I
1	cost(x) = '0.00~1000.00'	==>	revenue(x) = '0.00~500.00'	28.45	40.4				
2	cost(x) = '0.00~1000.00'	==>	revenue(x) = '500.00~1000.00'	20.46	29.05				
3	cost(x) = '0.00~1000.00'	==>	order_qty(x) = '0.00~100.00'	59.17	84.04				
4	cost(x) = '0.00~1000.00'	==>	revenue(x) = '1000.00~1500.00'	10.45	14.84				
5	cost(x) = '0.00~1000.00'	==>	region(x) = 'United States'	22.56	32.04				
6	cost(x) = '1000.00~2000.00'	==>	order_qty(x) = '0.00~100.00'	12.91	69.34				
7	order_qty(x) = '0.00~100.00'	==>	revenue(x) = '0.00~500.00'	28.45	34.54				
8	order_qty(x) = '0.00~100.00'	==>	cost(x) = '1000.00~2000.00'	12.91	15.67				
9	order_qty(x) = '0.00~100.00'	==>	region(x) = 'United States'	25.9	31.45				
10	order_qty(x) = '0.00~100.00'	==>	cost(x) = '0.00~1000.00'	59.17	71.86				
11	order_qty(x) = '0.00~100.00'	==>	product_line(x) = 'Tents'	13.52	16.42				
12	order_qty(x) = '0.00~100.00'	==>	revenue(x) = '500.00~1000.00'	19.67	23.88				
13	product_line(x) = 'Tents'	==>	order_qty(x) = '0.00~100.00'	13.52	98.72				
14	region(x) = 'United States'	==>	order_qty(x) = '0.00~100.00'	25.9	81.94				
15	region(x) = 'United States'	==>	cost(x) = '0.00~1000.00'	22.56	71.39				
16	revenue(x) = '0.00~500.00'	==>	cost(x) = '0.00~1000.00'	28.45	100				
17	revenue(x) = '0.00~500.00'	==>	order_qty(x) = '0.00~100.00'	28.45	100				
18	revenue(x) = '1000.00~1500.00'	==>	cost(x) = '0.00~1000.00'	10.45	96.75				
19	revenue(x) = '500.00~1000.00'	==>	cost(x) = '0.00~1000.00'	20.46	100				
20	revenue(x) = '500.00~1000.00'	==>	order_qty(x) = '0.00~100.00'	19.67	96.14				
21									
22									
23	cost(x) = '0.00~1000.00'	==>	revenue(x) = '0.00~500.00' AND order_qty(x) = '0.00~100.00'	28.45	40.4				
24	cost(x) = '0.00~1000.00'	==>	revenue(x) = '0.00~500.00' AND order_qty(x) = '0.00~100.00'	28.45	40.4				
25	cost(x) = '0.00~1000.00'	==>	revenue(x) = '500.00~1000.00' AND order_qty(x) = '0.00~100.00'	19.67	27.93				
26	cost(x) = '0.00~1000.00'	==>	revenue(x) = '500.00~1000.00' AND order_qty(x) = '0.00~100.00'	19.67	27.93				
27	cost(x) = '0.00~1000.00' AND order_qty(x) = '0.00~100.00'	==>	revenue(x) = '500.00~1000.00'	19.67	33.23				



# Visualization of Association Rule Using Plane Graph



# Visualization of Association Rule Using Rule Graph



# Iceberg Queries

- **Iceberg query**: Compute aggregates over one or a set of attributes only for those whose aggregate values is above certain threshold
- Example:  

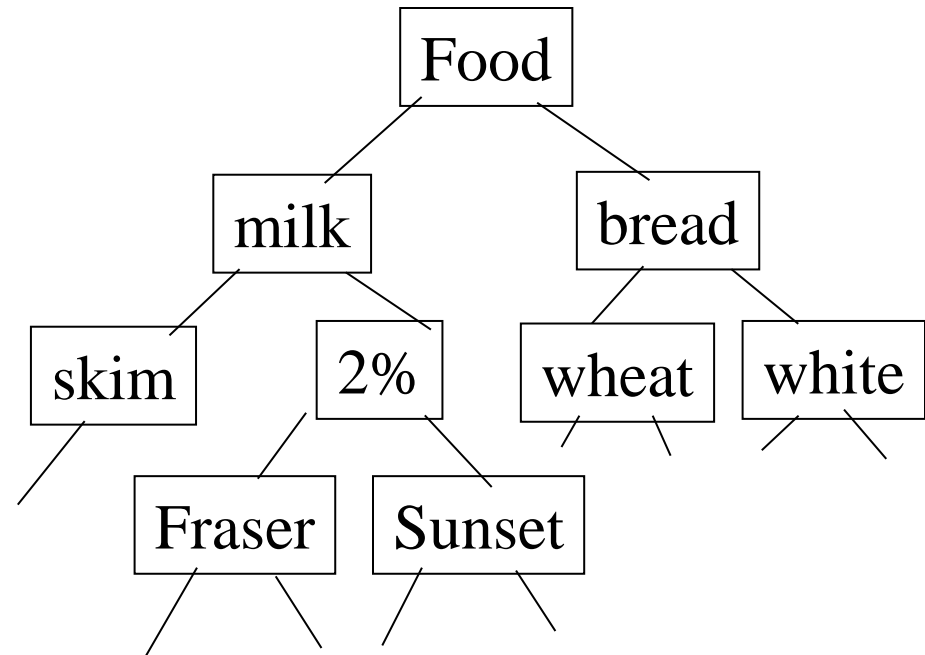
```
select P.custID, P.itemID, sum(P.qty)
from purchase P
group by P.custID, P.itemID
having sum(P.qty) >= 10
```
- **Compute** iceberg queries efficiently **by Apriori**:
  - First compute lower dimensions
  - Then compute higher dimensions only when **all** the lower ones are above the threshold

# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

# Multiple-Level Association Rules

- Items often form hierarchies.
- Items at the lower level are expected to have lower support.
- Rules regarding itemsets at appropriate levels could be quite useful.
- Transaction database can be encoded based on dimensions and levels
- We can explore shared multi-level mining



TID	Items
T1	{111, 121, 211, 221}
T2	{111, 211, 222, 323}
T3	{112, 122, 221, 411}
T4	{111, 121}
T5	{111, 122, 211, 221, 413}

# Mining Multi-Level Associations

- A top\_down, progressive deepening approach:
  - First find high-level strong rules:  
milk  $\rightarrow$  bread [20%, 60%].
  - Then find their lower-level “weaker” rules:  
2% milk  $\rightarrow$  wheat bread [6%, 50%].
- Variations at mining multiple-level association rules.
  - Level-crossed association rules:  
2% *milk*  $\rightarrow$  *Wonder* wheat bread
  - Association rules with multiple, alternative hierarchies:  
2% *milk*  $\rightarrow$  *Wonder* bread

# Multi-level Association: Uniform Support vs. Reduced Support

- **Uniform Support:** the same minimum support for all levels
  - + One minimum support threshold. No need to examine itemsets containing any item whose ancestors do not have minimum support.
  - — Lower level items do not occur as frequently. If support threshold
    - too high  $\Rightarrow$  miss low level associations
    - too low  $\Rightarrow$  generate too many high level associations
- **Reduced Support:** reduced minimum support at lower levels
  - There are 4 search strategies:
    - Level-by-level independent
    - Level-cross filtering by k-itemset
    - Level-cross filtering by single item
    - Controlled level-cross filtering by single item (level passage threshold)


# Uniform Support

## Multi-level mining with uniform support

**Level 1**  
**min\_sup = 5%**



**Level 2**  
**min\_sup = 5%**



[Back](#)




# Reduced Support

## Multi-level mining with reduced support

**Level 1**  
**min\_sup = 5%**

**Level 2**  
**min\_sup = 3%**



[Back](#)

# Multi-level Association: Redundancy Filtering

- Some rules may be **redundant** due to “ancestor” relationships between items.
- Example
  - milk  $\Rightarrow$  wheat bread [support = 8%, confidence = 70%]
  - 2% milk  $\Rightarrow$  wheat bread [support = 2%, confidence = 72%]
- We say the first rule is an ancestor of the second rule.
- A rule is redundant if its support is close to the “expected” value, based on the rule’s ancestor.

# Multi-Level Mining: Progressive Deepening


- A top-down, progressive deepening approach:
  - First mine high-level frequent items:  
milk (15%), bread (10%)
  - Then mine their lower-level “weaker” frequent itemsets:  
2% milk (5%), wheat bread (4%)
- Different  $\text{min\_support}$  threshold across multi-levels lead to different algorithms:
  - If adopting the same *min\_support* across multi-levels then toss  $t$  if any of  $t$ 's ancestors is infrequent.
  - If adopting reduced *min\_support* at lower levels then examine only those descendents whose ancestor's support is frequent/non-negligible.

# Progressive Refinement of Data Mining Quality

- Why progressive refinement?
  - Mining operator can be expensive or cheap, fine or rough
  - Trade speed with quality: step-by-step refinement.
- Superset coverage property:
  - Preserve all the positive answers—allow a positive false test but not a false negative test.
- Two- or multi-step mining:
  - First apply rough/cheap operator (superset coverage)
  - Then apply expensive algorithm on a substantially reduced candidate set (Koperski & Han, SSD'95).

# Progressive Refinement Mining of Spatial Association Rules

- Hierarchy of spatial relationship:
  - “g\_close\_to”: near\_by, touch, intersect, contain, etc.
  - First search for rough relationship and then refine it.
- Two-step mining of spatial association:
  - Step 1: rough spatial computation (as a filter)
    - Using MBR or R-tree for rough estimation.
  - Step2: Detailed spatial algorithm (as refinement)
    - Apply only to those objects which have passed the rough spatial association test (no less than *min\_support*)



# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

# Multi-Dimensional Association: Concepts

- Single-dimensional rules:  
 $\text{buys}(X, \text{"milk"}) \Rightarrow \text{buys}(X, \text{"bread"})$
- Multi-dimensional rules:  $\geq 2$  dimensions or predicates
  - Inter-dimension association rules (*no repeated predicates*)  
 $\text{age}(X, \text{"19-25"}) \wedge \text{occupation}(X, \text{"student"}) \Rightarrow \text{buys}(X, \text{"coke"})$
  - hybrid-dimension association rules (*repeated predicates*)  
 $\text{age}(X, \text{"19-25"}) \wedge \text{buys}(X, \text{"popcorn"}) \Rightarrow \text{buys}(X, \text{"coke"})$
- Categorical Attributes
  - finite number of possible values, no ordering among values
- Quantitative Attributes
  - numeric, implicit ordering among values


# Techniques for Mining MD Associations

- Search for frequent  $k$ -predicate set:
  - Example: {age, occupation, buys} is a 3-predicate set.
  - Techniques can be categorized by how age is treated:
- 1. Using static discretization of quantitative attributes
  - Quantitative attributes are statically discretized by using predefined concept hierarchies.
- 2. Quantitative association rules
  - Quantitative attributes are dynamically discretized into “bins” based on the distribution of the data.
- 3. Distance-based association rules
  - This is a dynamic discretization process that considers the distance between data points.



# Static Discretization of Quantitative Attributes

- Discretized prior to mining using concept hierarchy.
- Numeric values are replaced by ranges.
- In relational database, finding all frequent  $k$ -predicate sets will require  $k$  or  $k+1$  table scans.
- Data cube is well suited for mining.
- The cells of an  $n$ -dimensional cuboid correspond to the predicate sets.
- Mining from data cubes can be much faster.




# Quantitative Association Rules

- Numeric attributes are *dynamically* discretized
  - Such that the confidence or compactness of the rules mined is maximized.
- 2-D quantitative association rules:  $A_{\text{quan1}} \wedge A_{\text{quan2}} \Rightarrow A_{\text{cat}}$
- Cluster “adjacent”

association rules  
to form general  
rules using a 2-D  
grid.

- Example:

$\text{age}(X, "30-34") \wedge \text{income}(X, "24K - 48K")$   
 $\Rightarrow \text{buys}(X, "high\ resolution\ TV")$



# ARCS (Association Rule Clustering System)

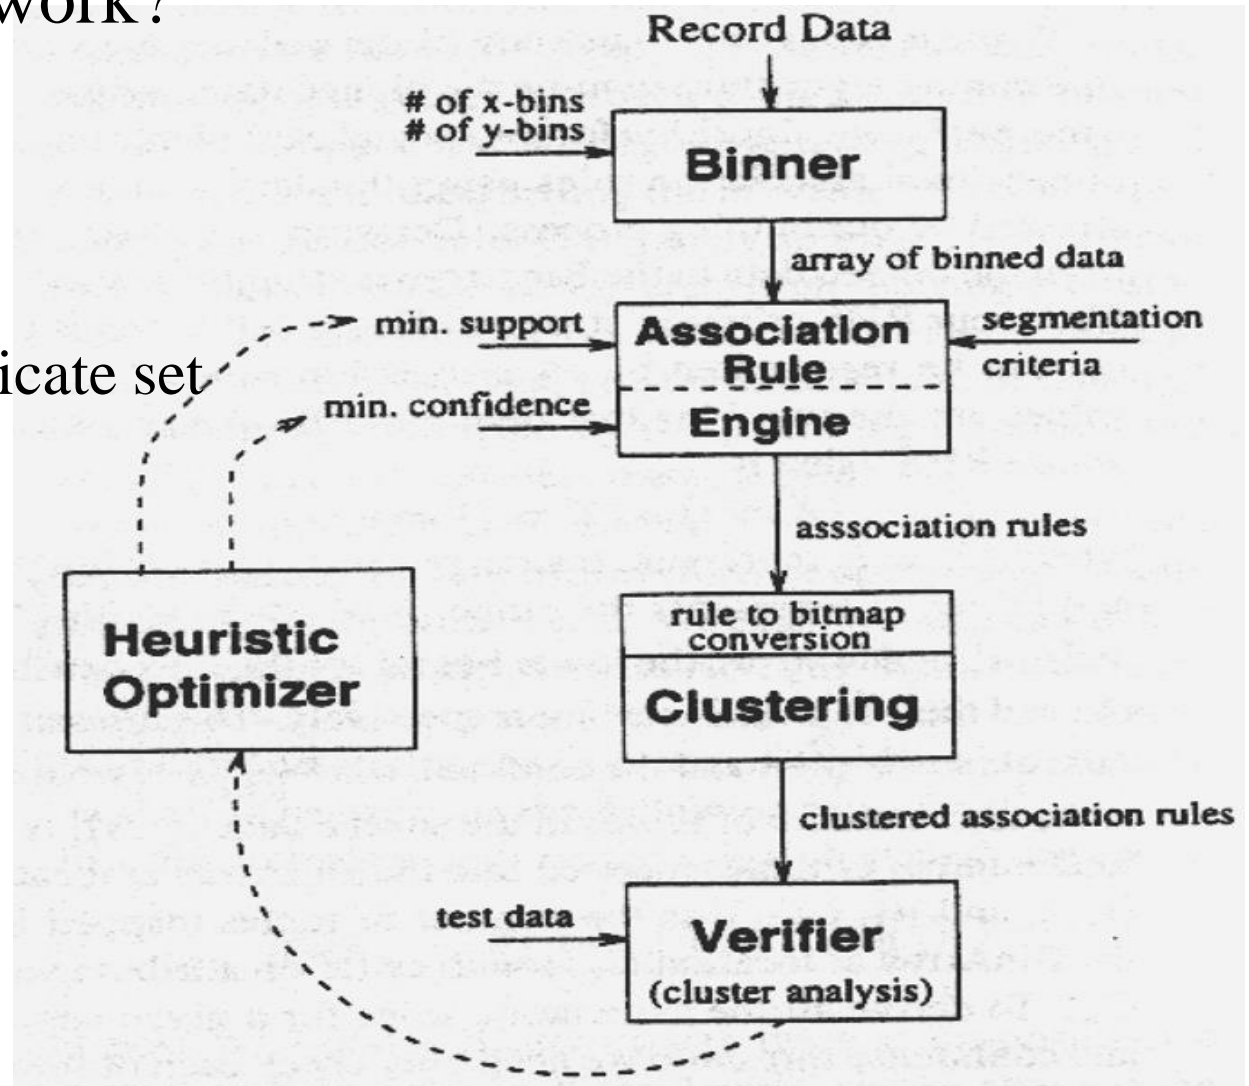
How does ARCS work?

1. Binning

2. Find frequent predicate set

3. Clustering

4. Optimize



# Limitations of ARCS

- Only quantitative attributes on LHS of rules.
- Only 2 attributes on LHS. (2D limitation)
- An alternative to ARCS
  - Non-grid-based
  - equi-depth binning
  - clustering based on a measure of *partial completeness* (*information lost due to partitioning*).
  - “***Mining Quantitative Association Rules in Large Relational Tables***” by R. Srikant and R. Agrawal.

# Mining Distance-based Association Rules

- Binning methods do not capture the semantics of interval data

Price(\$)	Equi-width (width \$10)	Equi-depth (depth 2)	Distance- based
7	[0,10]	[7,20]	[7,7]
20	[11,20]	[22,50]	[20,22]
22	[21,30]	[51,53]	[50,53]
50	[31,40]		
51	[41,50]		
53	[51,60]		

- Distance-based partitioning, more meaningful discretization considering:
  - density/number of points in an interval
  - “closeness” of points in an interval

# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

# Interestingness Measures

- Objective measures

Two popular measurements:

- *support*; and

- *confidence*

- Subjective measures (Silberschatz & Tuzhilin, KDD95)

A rule (pattern) is interesting if

- it is *unexpected* (surprising to the user); and/or

- *actionable* (the user can do something with it)

- From association to correlation and causal structure analysis.

- Association does not necessarily imply correlation or causal relationships

# Criticism to Support and Confidence

- Example 1: (Aggarwal & Yu, PODS98)
  - Among 5000 students
    - 3000 play basketball
    - 3750 eat cereal
    - 2000 both play basket ball and eat cereal
  - *play basketball*  $\Rightarrow$  *eat cereal* [40%, 66.7%] is misleading because the overall percentage of students eating cereal is 75% which is higher than 66.7%.
  - *play basketball*  $\Rightarrow$  *not eat cereal* [20%, 33.3%] is far more accurate, although with lower support and confidence

	basketball	not basketball	sum(row)
cereal	2000	1750	3750
not cereal	1000	250	1250
sum(col.)	3000	2000	5000



# Criticism to Support and Confidence

- Example 2:

- X and Y: positively correlated,
- X and Z, negatively related
- support and confidence of  $X \Rightarrow Z$  dominates

X	1	1	1	1	0	0	0	0
Y	1	1	0	0	0	0	0	0
Z	0	1	1	1	1	1	1	1

- We need a measure of dependent or correlated events

$$corr_{A,B} = \frac{P(A \cup B)}{P(A)P(B)}$$

Rule	Support	Confidence
$X \Rightarrow Y$	25%	50%
$X \Rightarrow Z$	37.50%	75%

- ...= $P(B|A)/P(B)$  is also called the **lift** of rule  $A \Rightarrow B$

# Other Interestingness Measures: Interest

- Interest (correlation, lift)  $\frac{P(A \wedge B)}{P(A)P(B)}$ 
  - taking both  $P(A)$  and  $P(B)$  in consideration
  - $P(A \wedge B) = P(B) * P(A)$ , if A and B are independent events
  - A and B negatively correlated, if the value is less than 1;  
otherwise A and B positively correlated

X	1	1	1	1	0	0	0	0
Y	1	1	0	0	0	0	0	0
Z	0	1	1	1	1	1	1	1

Itemset	Support	Interest
X,Y	25%	2
X,Z	37.50%	0.9
Y,Z	12.50%	0.57

# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

# Constraint-Based Mining

- Interactive, exploratory mining giga-bytes of data?
  - Could it be real? — Making good use of constraints!
- What kinds of constraints can be used in mining?
  - **Knowledge type constraint**: classification, association, etc.
  - **Data constraint**: SQL-like queries
    - Find product pairs sold together in **Vancouver** in **Dec.**'98.
  - **Dimension/level constraints**:
    - in relevance to **region, price, brand, customer category**.
  - **Rule constraints**
    - On the form of the rules to be mined (e.g., # of predicates, etc)
    - small sales (price < \$10) triggers big sales (sum > \$200).
  - **Interestingness constraints**:
    - Thresholds on measures of interestingness
    - strong rules (min\_support  $\geq$  3%, min\_confidence  $\geq$  60%).

# Rule Constraints in Association Mining

- Two kind of rule constraints:
  - Rule form constraints: meta-rule guided mining.
    - $P(x, y) \wedge Q(x, w) \rightarrow \text{takes}(x, \text{“database systems”})$ .
  - Rule (content) constraint: constraint-based query optimization (Ng, et al., SIGMOD'98).
    - $\text{sum}(\text{LHS}) < 100 \wedge \text{min}(\text{LHS}) > 20 \wedge \text{count}(\text{LHS}) > 3 \wedge \text{sum}(\text{RHS}) > 1000$
- 1-variable vs. 2-variable constraints (Lakshmanan, et al. SIGMOD'99):
  - 1-var: A constraint confining only one side (L/R) of the rule, e.g., as shown above.
  - 2-var: A constraint confining both sides (L and R).
    - $\text{sum}(\text{LHS}) < \text{min}(\text{RHS}) \wedge \text{max}(\text{RHS}) < 5 * \text{sum}(\text{LHS})$

# Constraint-Based Association Query

- Database: (1) trans (TID, Itemset ), (2) itemInfo (Item, Type, Price)
- A constrained assoc. query (CAQ) is in the form of  $\{(S_1, S_2)/C\}$ ,
  - where C is a set of constraints on  $S_1, S_2$  including frequency constraint
- A classification of (single-variable) constraints:
  - Class constraint:  $S \subset A$ . *e.g.  $S \subset Item$*
  - Domain constraint:
    - $S \theta v, \theta \in \{=, \neq, <, \leq, >, \geq\}$ . *e.g.  $S.Price < 100$*
    - $v \theta S, \theta$  is  $\in$  or  $\notin$ . *e.g.  $snacks \notin S.Type$*
    - $V \theta S$ , or  $S \theta V, \theta \in \{\subseteq, \subset, \not\subset, =, \neq\}$ 
      - *e.g.  $\{snacks, sodas\} \subseteq S.Type$*
  - Aggregation constraint:  $agg(S) \theta v$ , where  $agg$  is in  $\{min, max, sum, count, avg\}$ , and  $\theta \in \{=, \neq, <, \leq, >, \geq\}$ .
    - *e.g.  $count(S_1.Type) = 1, avg(S_2.Price) < 100$*

# Constrained Association Query Optimization Problem

- Given a CAQ =  $\{ (S_1, S_2) / C \}$ , the algorithm should be :
  - **sound**: It only finds frequent sets that satisfy the given constraints C
  - **complete**: All frequent sets satisfy the given constraints C are found
- A naïve solution:
  - Apply Apriori for finding all frequent sets, and **then** to test them for constraint satisfaction one by one.
- Other approach:
  - Comprehensive analysis of the properties of constraints and try to **push them as deeply as possible inside** the frequent set computation.

# Anti-monotone and Monotone Constraints

- A constraint  $C_a$  is **anti-monotone** iff. for any pattern  $S$  not satisfying  $C_a$ , none of the super-patterns of  $S$  can satisfy  $C_a$
- A constraint  $C_m$  is **monotone** iff. for any pattern  $S$  satisfying  $C_m$ , every super-pattern of  $S$  also satisfies it



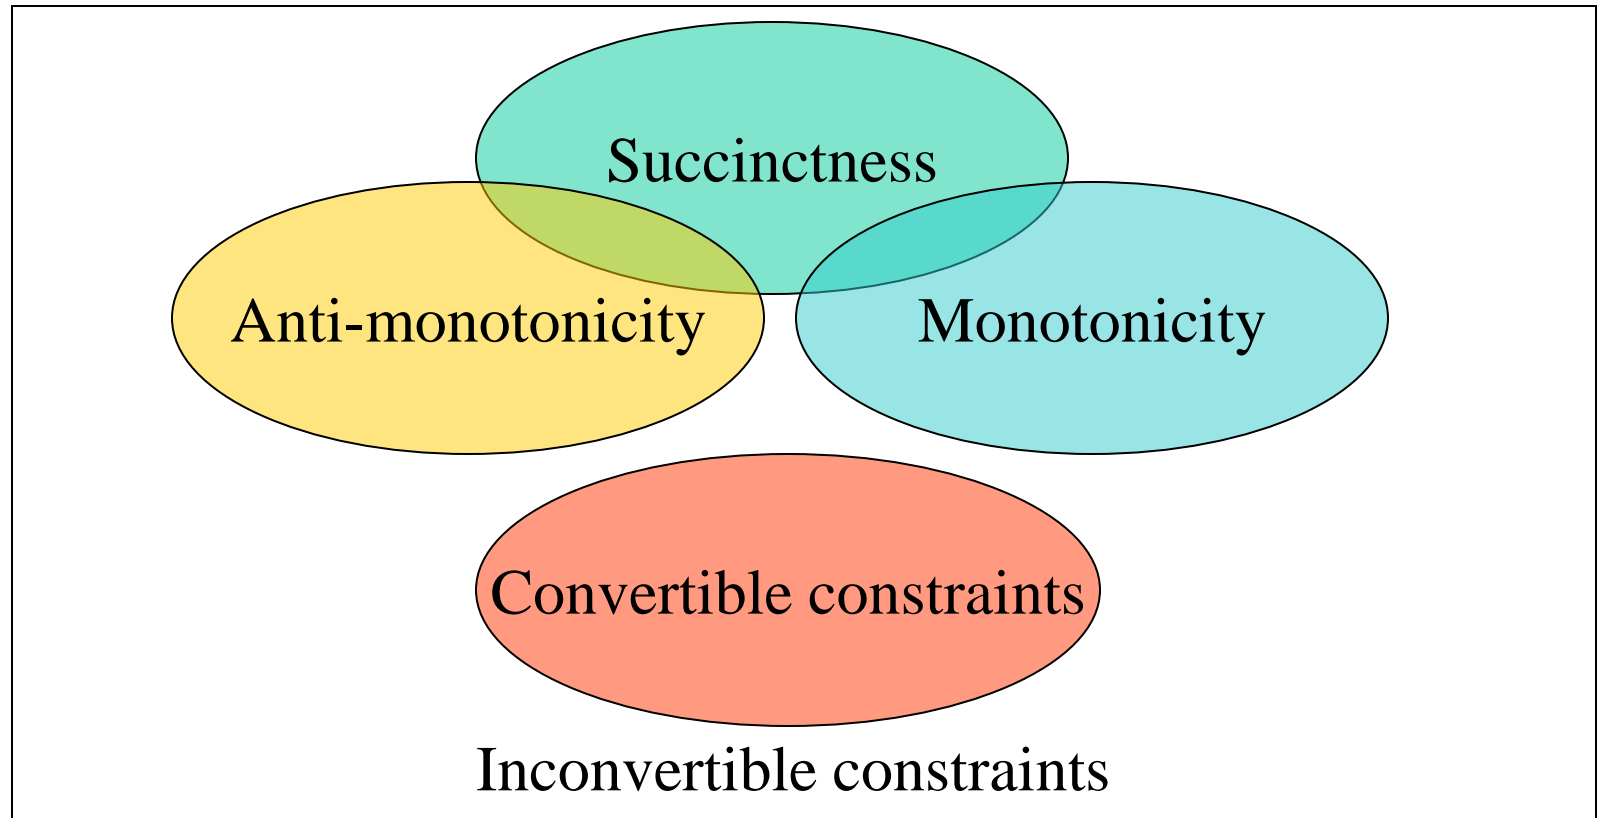
# Succinct Constraint

- A subset of item  $I_s$  is a **succinct set**, if it can be expressed as  $\sigma_p(I)$  for some selection predicate  $p$ , where  $\sigma$  is a selection operator
- $SP \subseteq 2^I$  is a succinct **power set**, if there is a fixed number of succinct set  $I_1, \dots, I_k \subseteq I$ , s.t.  $SP$  can be expressed in terms of the strict power sets of  $I_1, \dots, I_k$  using union and minus
- A constraint  $C_s$  is **succinct** provided  $SAT_{C_s}(I)$  is a succinct power set

# Convertible Constraint

- Suppose all items in patterns are listed in a total order  $R$
- A constraint  $C$  is **convertible anti-monotone** iff a pattern  $S$  satisfying the constraint implies that each suffix of  $S$  w.r.t.  $R$  also satisfies  $C$
- A constraint  $C$  is **convertible monotone** iff a pattern  $S$  satisfying the constraint implies that each pattern of which  $S$  is a suffix w.r.t.  $R$  also satisfies  $C$

# Relationships Among Categories of Constraints



# Property of Constraints: Anti-Monotone

- Anti-monotonicity: *If a set  $S$  violates the constraint, any superset of  $S$  violates the constraint.*
- Examples:
  - $\text{sum}(S.\text{Price}) \leq v$  is anti-monotone
  - $\text{sum}(S.\text{Price}) \geq v$  is not anti-monotone
  - $\text{sum}(S.\text{Price}) = v$  is partly anti-monotone
- Application:
  - Push “ $\text{sum}(S.\text{price}) \leq 1000$ ” deeply into iterative frequent set computation.

# Characterization of Anti-Monotonicity Constraints

$S \theta v, \theta \in \{=, \leq, \geq\}$	<b>yes</b>
$v \in S$	<b>no</b>
$S \supseteq V$	<b>no</b>
$S \subseteq V$	<b>yes</b>
$S = V$	<b>partly</b>
$\min(S) \leq v$	<b>no</b>
$\min(S) \geq v$	<b>yes</b>
$\min(S) = v$	<b>partly</b>
$\max(S) \leq v$	<b>yes</b>
$\max(S) \geq v$	<b>no</b>
$\max(S) = v$	<b>partly</b>
$\text{count}(S) \leq v$	<b>yes</b>
$\text{count}(S) \geq v$	<b>no</b>
$\text{count}(S) = v$	<b>partly</b>
$\text{sum}(S) \leq v$	<b>yes</b>
$\text{sum}(S) \geq v$	<b>no</b>
$\text{sum}(S) = v$	<b>partly</b>
$\text{avg}(S) \theta v, \theta \in \{=, \leq, \geq\}$	<b>convertible</b>
(frequent constraint)	(yes)

# Example of Convertible Constraints: $\text{Avg}(S) \geq V$

- Let  $R$  be the value descending order over the set of items
  - E.g.  $I = \{9, 8, 6, 4, 3, 1\}$
- $\text{Avg}(S) \geq v$  is convertible monotone w.r.t.  $R$ 
  - If  $S$  is a suffix of  $S_1$ ,  $\text{avg}(S_1) \geq \text{avg}(S)$ 
    - $\{8, 4, 3\}$  is a suffix of  $\{9, 8, 4, 3\}$
    - $\text{avg}(\{9, 8, 4, 3\}) = 6 \geq \text{avg}(\{8, 4, 3\}) = 5$
  - If  $S$  satisfies  $\text{avg}(S) \geq v$ , so does  $S_1$ 
    - $\{8, 4, 3\}$  satisfies constraint  $\text{avg}(S) \geq 4$ , so does  $\{9, 8, 4, 3\}$

# Property of Constraints: Succinctness

- Succinctness:
  - For any set  $S_1$  and  $S_2$  satisfying  $C$ ,  $S_1 \cup S_2$  satisfies  $C$
  - Given  $A_1$  is the sets of size 1 satisfying  $C$ , then any set  $S$  satisfying  $C$  are based on  $A_1$ , i.e., it contains a subset belongs to  $A_1$ ,
- Example :
  - $\text{sum}(S.\text{Price}) \geq v$  is not succinct
  - $\text{min}(S.\text{Price}) \leq v$  is succinct
- Optimization:
  - If  $C$  is succinct, then  $C$  is pre-counting prunable. The satisfaction of the constraint alone is not affected by the iterative support counting.

# Characterization of Constraints by Succinctness

$S \theta v, \theta \in \{=, \leq, \geq\}$	<b>Yes</b>
$v \in S$	<b>yes</b>
$S \supseteq V$	<b>yes</b>
$S \subseteq V$	<b>yes</b>
$S = V$	<b>yes</b>
$\min(S) \leq v$	<b>yes</b>
$\min(S) \geq v$	<b>yes</b>
$\min(S) = v$	<b>yes</b>
$\max(S) \leq v$	<b>yes</b>
$\max(S) \geq v$	<b>yes</b>
$\max(S) = v$	<b>yes</b>
$\text{count}(S) \leq v$	<b>weakly</b>
$\text{count}(S) \geq v$	<b>weakly</b>
$\text{count}(S) = v$	<b>weakly</b>
$\text{sum}(S) \leq v$	<b>no</b>
$\text{sum}(S) \geq v$	<b>no</b>
$\text{sum}(S) = v$	<b>no</b>
$\text{avg}(S) \theta v, \theta \in \{=, \leq, \geq\}$	<b>no</b>
(frequent constraint)	<b>(no)</b>



# Agenda

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Mining multilevel association rules from transactional databases
- Mining multidimensional association rules from transactional databases and data warehouse
- From association mining to correlation analysis
- Constraint-based association mining
- Summary

# Summary

- Association rule mining
  - probably the most significant contribution from the database community in KDD
  - large number of papers
- Many interesting issues have been explored
- An interesting research direction
  - Association analysis in other types of data: spatial data, multimedia data, time series data, etc.